

CSO Assignment - 2

Task 0 : System Information

The complete information regarding the system can be found in `task0.txt` within the directory.

Task 1: Optimising Matrix Multiplication

The original algorithm takes **246.048780163 seconds** with `-O2` optimisation flag and the `cache-misses` for this algorithm is **4,48,28,26,481 misses**, as given below.

```
~/CSO/Assignment 2 took 3s
→ sudo perf stat -d -e cache-misses ./a.out
Enter number of rows and columns of first matrix
Enter elements of first matrix
Enter number of rows and columns of second matrix
Enter elements of second matrix
Product of the matrices:
Required Time: 245.729805
Performance counter stats for './a.out':

   4,48,28,26,481      cache-misses                  (79.99%)
   86,39,06,97,040     L1-dcache-loads                  (80.01%)
   64,05,11,43,101     L1-dcache-load-misses           # 74.14% of all L1-dcache hits (80.00%)
   41,71,80,54,753     LLC-loads                       (80.00%)
    1,45,02,25,374     LLC-load-misses                 # 3.48% of all LL-cache hits  (80.00%)

   246.048780163 seconds time elapsed

   245.570546000 seconds user
    0.168001000 seconds sys
```

The algorithm given can be optimised in several ways to reduce the amount of cache misses and increase performance. For instance the order of the access needs to be proper, ie: for matrix multiplication for matrix $A[i][k]$ and $B[k][j]$ the order of loop should be i , j and k .

Another possible optimisation is a simple tiling algorithm which involves dividing the matrix multiplication to smaller blocks. This means that the block in consideration will not exceed the cache size and hence will improve the performance by a great margin. The block-size I have taken here is `16`.

```
10 void matrixMul(int m, int n, int q)
11 {
12     int sum;
13     int i, j, k, jj, kk;
14
15     for (kk = 0; kk < n; kk += BLOCK)
16         for (jj = 0; jj < q; jj += BLOCK)
17             for (i = 0; i < m; i++)
18                 for (k = kk; k < ((n < kk + BLOCK) ? n : kk + BLOCK); k++)
19                     for (j = jj; j < ((q < jj + BLOCK) ? q : jj + BLOCK); j++)
20                         M[i][j] += A[i][k] * B[k][j];
21 }
22 }
```

The modified algorithm takes **37.812287392 seconds** with `-O2` optimisation flag and the `cache-misses` for it is **1,59,59,89,815 misses** as given below.

```
~/CS0/2019113022
→ sudo perf stat -d -e cache-misses ./a.out
Required Time: 37.600704
Performance counter stats for './a.out':

1,59,59,89,815      cache-misses              (79.96%)
1,28,69,26,73,889    L1-dcache-loads           (80.04%)
1,12,83,84,263      L1-dcache-load-misses     #    0.88% of all L1-dcache hits (79.99%)
46,59,73,012        LLC-loads                 (80.03%)
9,74,81,025         LLC-load-misses          #   20.92% of all LL-cache hits (79.98%)

37.812287392 seconds time elapsed

37.561075000 seconds user
0.051990000 seconds sys
```

Task 2: Optimising Merge Sort

The algorithm given is originally is Merge Sort, which is a very optimal sorting algorithm, with the best possible sorting complexity of $O(n \log n)$. However, this algorithm is not very good for very small number of elements.

The original algorithm takes **0.844581871 seconds** with `-O2` optimisation flag and the `cache-misses` for this algorithm is **1,34,94,479 misses**, as given below.

```
~/CS0/Assignment 2
→ sudo perf stat -d -e cache-misses ./a.out
Sorted array is

Performance counter stats for './a.out':

1,34,94,479      cache-misses              (79.95%)
56,82,91,997     L1-dcache-loads           (80.13%)
1,80,93,597      L1-dcache-load-misses     #    3.18% of all L1-dcache hits (80.08%)
3,11,033         LLC-loads                 (80.08%)
1,27,909         LLC-load-misses          #   41.12% of all LL-cache hits (79.76%)

0.844581871 seconds time elapsed

0.832116000 seconds user
0.012059000 seconds sys
```

The work around is a sort called Tim Sort. This is a form of tiling sort, in which the array is divided into blocks, and sorted according with insertions sort. Then these arrays are then merged, similar to merge sort.

```

71 void timSort(int arr[], int n)
72 {
73     int i;
74
75     for (i = 0; i < n; i += BLOCK)
76         insertionSort(arr, i, min((i + BLOCK - 1), (n - 1)));
77
78     for (i = BLOCK; i < n; i = 2*i)
79     {
80         for (int l = 0; l < n; l += 2*i)
81         {
82             int m = l + i - 1, r = min((l + 2*i - 1), (n - i));
83
84             if(m < l)
85                 mergeSort(arr, l, m, r);
86         }
87     }
88 }

```

The modified algorithm takes **0.137804001 seconds** with **-O2** optimisation flag and the **cache-misses** for it is **12,59,108 misses** as given below.

```

~/CS0/2019113022$ sudo perf stat -d -e cache-misses ./a.out
Required Time: 0.135333
Performance counter stats for './a.out':
 12,59,108      cache-misses
15,25,20,063    L1-dcache-loads
10,03,015      L1-dcache-load-misses
 41,884        LLC-loads
 26,980        LLC-load-misses

0.137804001 seconds time elapsed

0.129705000 seconds user
0.008106000 seconds sys

```